



Course material at: <https://git.io/fjFga>

Regression, Overfitting and Optimization


Monday, Lecture 2
FASE ML Bootcamp

Based on material from Joseph Redmon, Christopher Bishop,
Polo Chau
Elements of Statistical Learning by Hastie, Tibshirani, Friedman



Regression


Data points (x, t) generated by adding noise to $\sin(2\pi x)$







Loss Function!

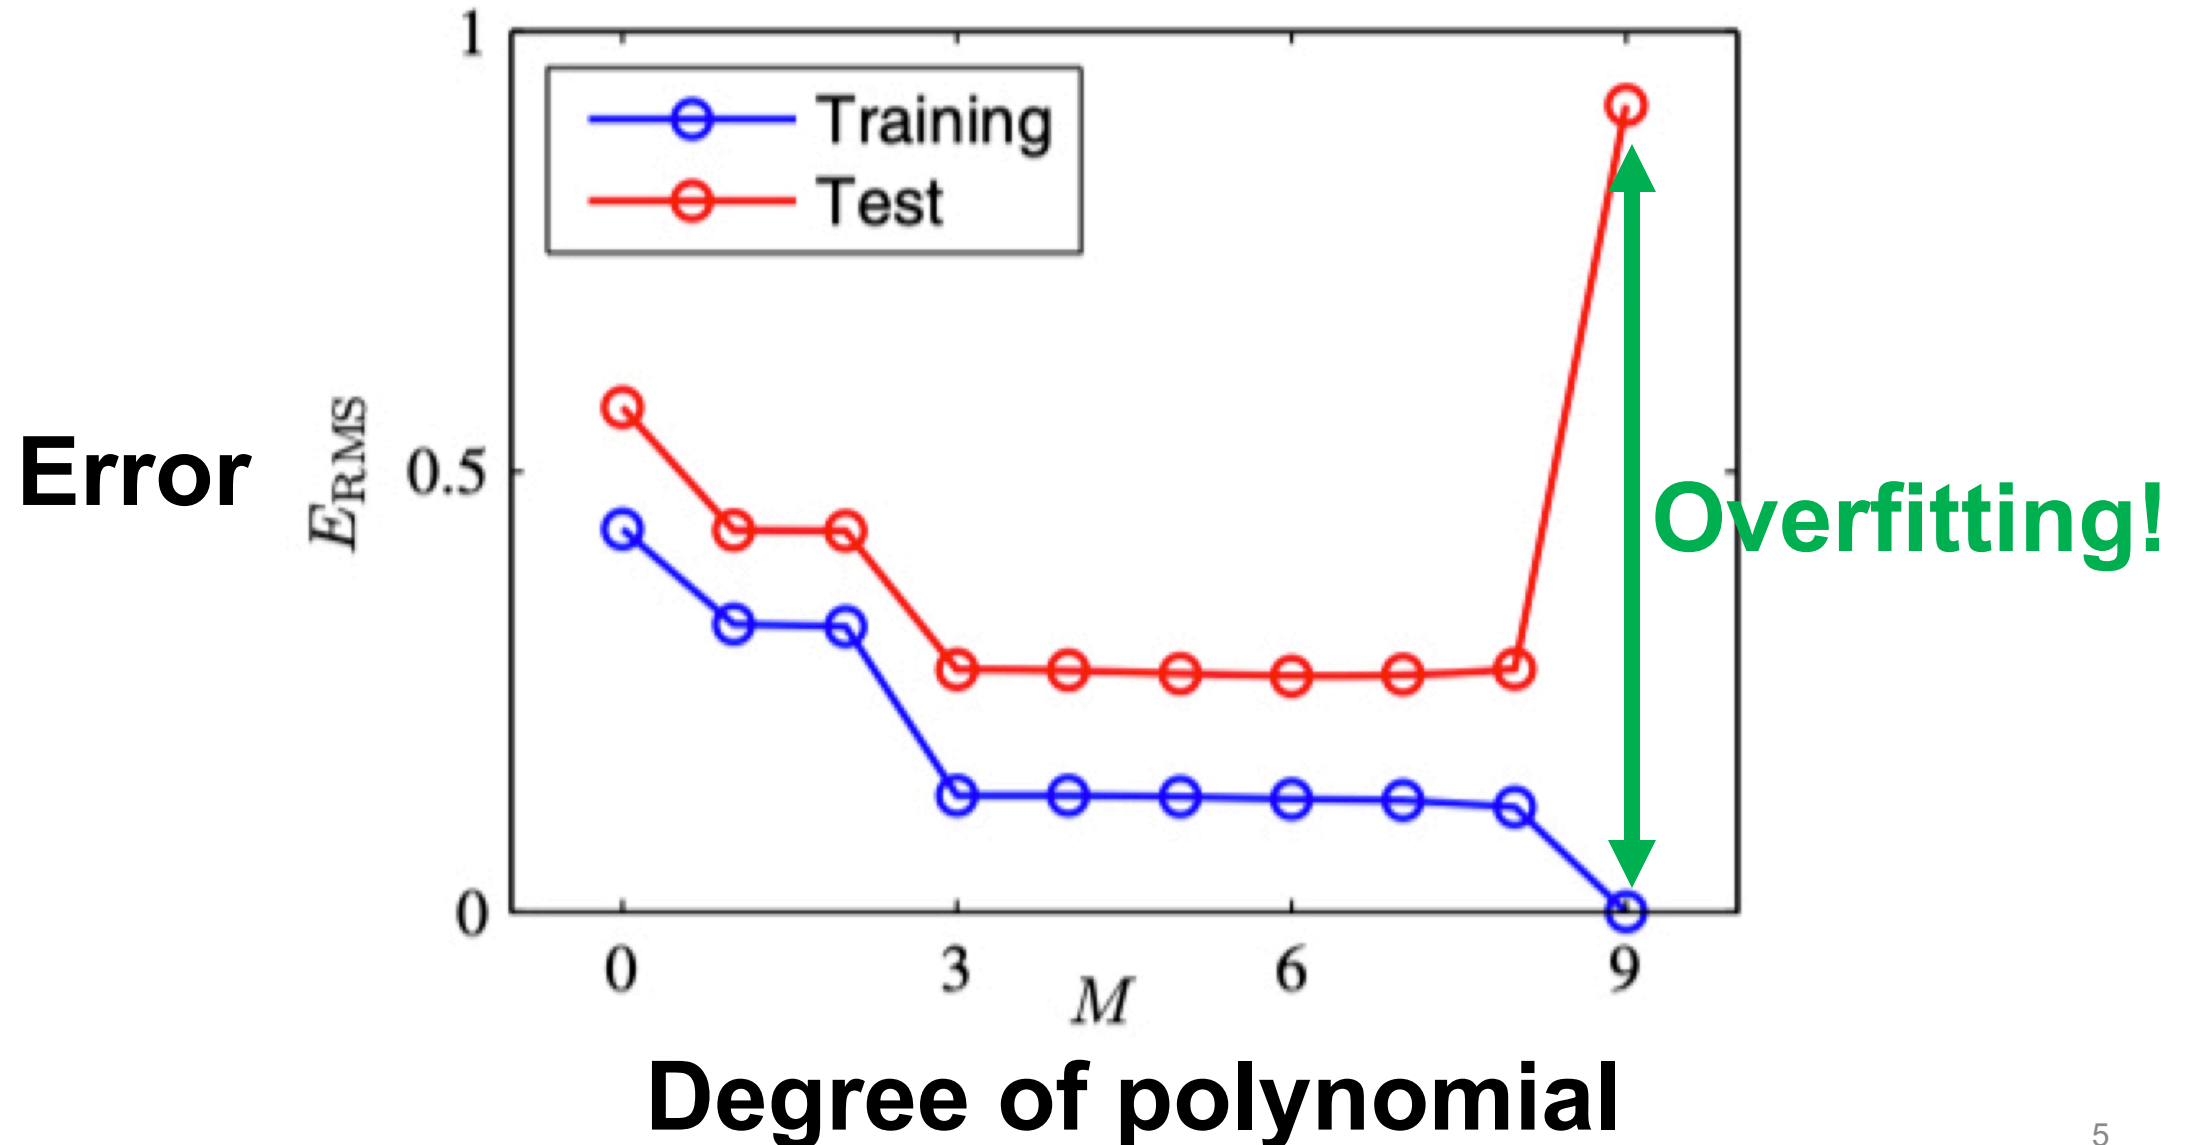
Errors in Regression

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$






Error as a function of degree of polynomial



Dealing with Overfitting: (1) more data!



Dealing with Overfitting: (2) regularization

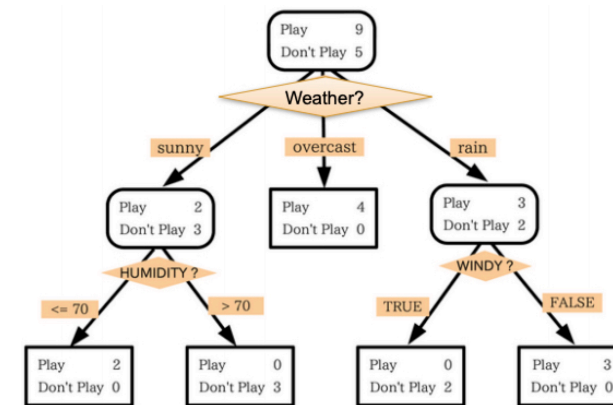
minimize:

$$\tilde{E}(\mathbf{w}) = \underbrace{\frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2}_{\text{Model's squared error}} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{Penalty on parameters}}$$

Regularization in Classification

- Decision Trees

- Limit the **depth** of the tree
- Prune subtrees after training



- Support Vector Machines

- Built-in, tunable **regularization term**


- Logistic Regression

- Can add a tunable **regularization term** to MLE objective


Penalize large
parameter
values!

Underfitting: model is too simplistic!

Training loss too high


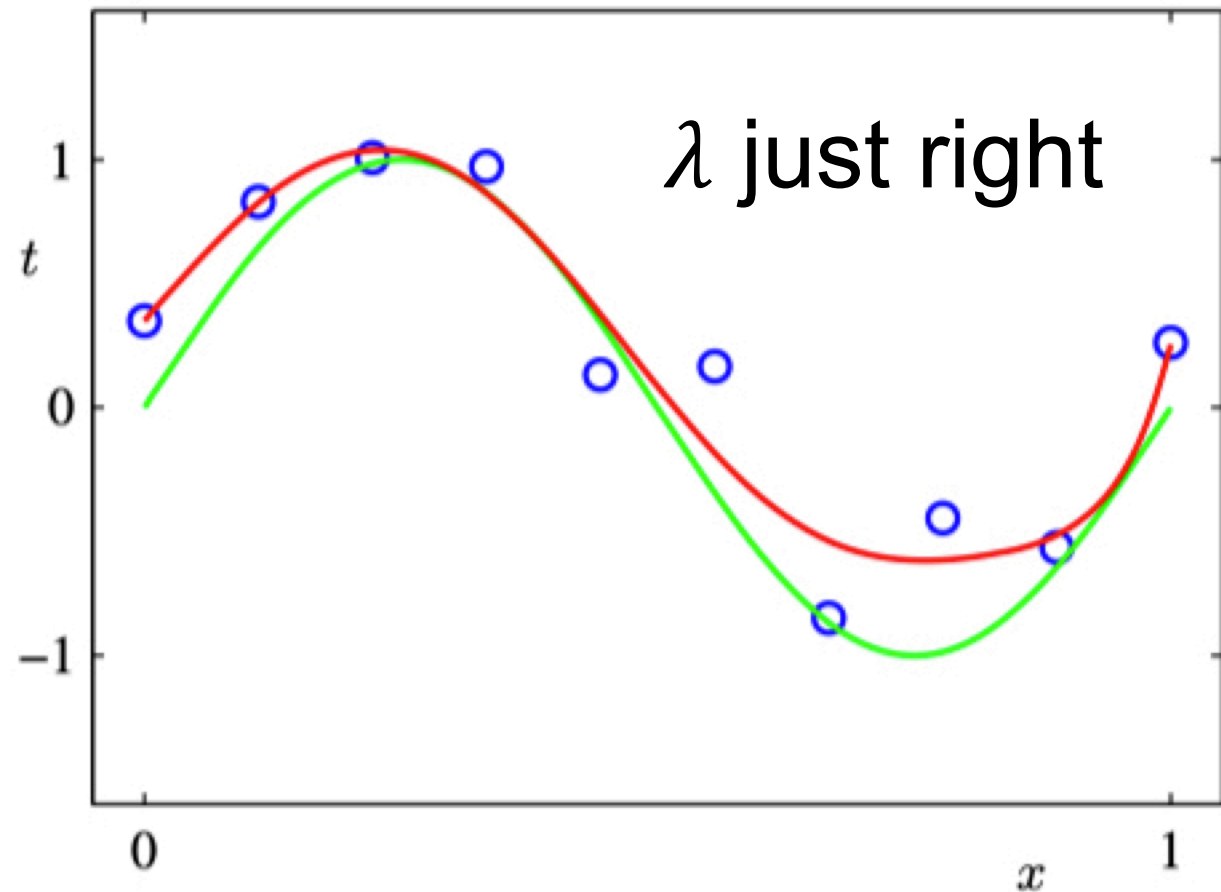


Use a more complex model



Dealing with Overfitting: (2) regularization

minimize: $\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$



Parameters vs **Hyper**-parameters

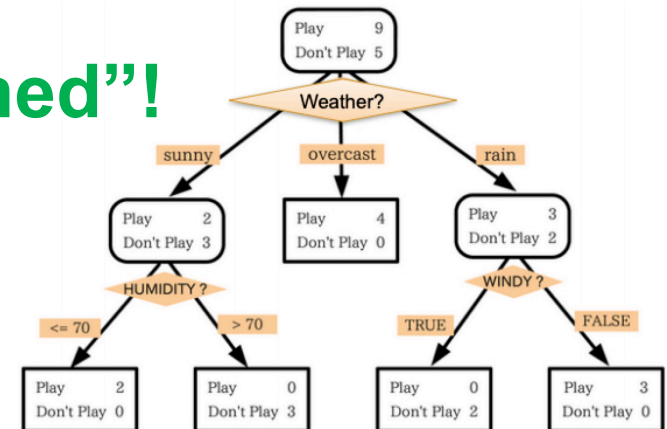
Example hyper-parameters

- Decision tree: maximum depth of the tree
- Polynomial regression: regularization coefficient $\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$
- Hyper-parameters are **determined through trial-and-error / cross validation**

Example parameters: **directly optimized, “learned”!**

- Decision tree: attributes you split on
- Logistic regression: weights β

$$p(x; \beta) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_d x_d)}}$$



Regression

Examples:

- Stock price prediction
- Forecasting epidemics
- Weather prediction


We will look at:

- Linear Regression
- Ridge Regression
- LASSO





Regression

What is the temperature going to be tomorrow?



https://medium.com/@ali_88273/regression-vs-classification-87c224350d69



Linear Regression

Data: $S = \{(x_i, y_i)\}_{i=1, \dots, n}$

x_i : data example with **d** attributes


y_i : target of example (what you care about)

Model:

$$f(x; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \dots + \beta_d x_d$$

Loss function: Residual Sum of Squares

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(x_i; \boldsymbol{\beta}))^2$$



Linear Regression

Minimizing RSS to find β^* :


- Closed-form solution!
- If closed-form fails, can use standard optimization methods

Model:

$$f(x; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$

Loss function: Residual Sum of Squares

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(x_i; \boldsymbol{\beta}))^2$$



Ridge Regression

- Linear Regression uses all features; model may be complicated
- **Ridge Regression** penalizes large parameter values

Model:

$$f(x; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$

Loss function: Residual Sum of Squares + **penalty** term

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(x_i; \boldsymbol{\beta}))^2 + \lambda \sum_{j=0}^d \beta_j^2$$

Lasso Regression

- As in Ridge Regression, Lasso penalizes large parameters
- Penalizes **absolute** instead of squared coefficient values
- **Zeroes out** more coefficients **BUT** optimization is more involved

Model:

$$f(x; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$

Loss function: Residual Sum of Squares + **penalty** term

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(x_i; \boldsymbol{\beta}))^2 + \lambda \sum_{j=0}^d |\beta_j|$$

Example: Prostate Cancer

Stamey et al. (1989)


- x: cancer volume, prostate weight, age, ...
- y: amount of prostate-specific antigen

Term	LR	Best Subset	Ridge	Lasso
Intercept	2.465	2.477	2.452	2.468
lcavol	0.680	0.740	0.420	0.533
lweight	0.263	0.316	0.238	0.169
age	-0.141		-0.046	
lbph	0.210		0.162	0.002
svi	0.305		0.227	0.094
lcp	-0.288		0.000	
gleason	-0.021		0.040	
pgg45	0.267		0.133	
Test Error	0.521	0.492	0.492	0.479
Std Error	0.179	0.143	0.165	0.164

Bias / Variance tradeoff

Based on Joseph Redmon's slides,
UW Computer Vision class

- Bias
 - Error from assumptions model makes about data
 - Linear model assumes data is linear, bad for data that isn't




Bias / Variance tradeoff

Based on Joseph Redmon's slides,
UW Computer Vision class

- Bias


- Error from assumptions model makes about data
- Linear model assumes data is linear, bad for data that isn't
- Quadratic is better



Bias / Variance tradeoff

Based on Joseph Redmon's slides,
UW Computer Vision class



- Variance
 - Algorithm's sensitivity to noise
 - More complex algorithms are more sensitive!



Bias / Variance tradeoff

Based on Joseph Redmon's slides,
UW Computer Vision class

- Variance
 - Algorithm's sensitivity to noise
 - More complex algorithms are more sensitive!





Bias / Variance tradeoff

Based on Joseph Redmon's slides,
UW Computer Vision class

- Variance

- Algorithm's sensitivity to noise
- More complex algorithms are more sensitive!
- High variance hurts generalization, *overfitting*




$$\text{Error} = \text{Noise} + \text{Bias} + \text{Variance}$$

- Noise
 - Random variations in data
- Bias
 - **Error from assumptions** model makes about data
 - Less complex algorithms -> more assumptions about data
- Variance
 - Algorithm's **sensitivity to noise**
 - More complex algorithms are more sensitive!
 - High variance hurts generalization

Linear regression

Based on Joseph Redmon's slides,
UW Computer Vision class


- $f^*(x) = ax + b$
- **High bias:** linear assumption
- **Low variance**
- **Benefits:**
 - Closed form solution
 - Fast to compute for new data
- **Weaknesses:**
 - Not very powerful, assumes linear
 - Underfit more interesting data



Nearest neighbor regression

Based on Joseph Redmon's slides,
UW Computer Vision class


- $f^*(x) = f(x')$ for nearest x' in training set
- **Low bias**: no assumptions about data
- **High variance**: very sensitive to training set
- Benefits:
 - Super easy to implement
 - Easy to understand
 - Arbitrarily powerful, esp with lots of data
- Weaknesses:
 - Hard to scale
 - Prone to **overfitting** to noise



ML as an Optimization problem: Regression

minimize:

$$\tilde{E}(\mathbf{w}) = \underbrace{\frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2}_{\text{Model's squared error}} + \underbrace{\frac{\lambda}{2} \|\mathbf{w}\|^2}_{\text{Penalty on parameters}}$$



Model's squared error

Penalty on parameters

ML as an Optimization problem: Logistic Regression

$$p(x; \boldsymbol{\beta}) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_d x_d)}}$$

Data: $S = \{(x_i, y_i)\}_{i=1, \dots, n}$ | x_i : example with d attributes (age, #pregnancies, ...)
 y_i : cervical cancer diagnosis (0 or 1)

Maximum Likelihood Estimation (MLE)

Likelihood of observing the data for a given $\boldsymbol{\beta}$:

$$\prod_{i=1}^n p(x_i; \boldsymbol{\beta})^{y_i} \times (1 - p(x_i; \boldsymbol{\beta}))^{1-y_i}$$


MLE seeks parameters $\boldsymbol{\beta}$ that maximize the likelihood

The optimal parameters, $\boldsymbol{\beta}^*$, can be found by optimization

Convexity


- Graphically... which of these is convex?

For convex problems, gradient descent gets us to **global minima**. It's that easy!



Non-convexity

Usually no easy way to find global or local optima, harder to optimize




Gradient descent

Based on Joseph Redmon's slides,
UW Computer Vision class

For some loss function $L(\mathbf{w})$, gradient $\nabla L(\mathbf{w})$ points towards in direction of steepest ascent.

In 1d, either points left or right




Gradient descent

Based on Joseph Redmon's slides,
UW Computer Vision class

For some loss function $L(\mathbf{w})$, gradient $\nabla L(\mathbf{w})$ points towards in direction of steepest ascent.

In 1d, either points left or right




Gradient descent

Based on Joseph Redmon's slides,
UW Computer Vision class

For some loss function $L(\mathbf{w})$, gradient $\nabla L(\mathbf{w})$ points towards in direction of steepest ascent.

In 1d, either points left or right




Gradient descent

Based on Joseph Redmon's slides,
UW Computer Vision class

For some loss function $L(\mathbf{w})$, gradient $\nabla L(\mathbf{w})$ points towards in direction of steepest ascent.

In 1d, either points left or right



Gradient descent

Based on Joseph Redmon's slides,
UW Computer Vision class

For some loss function $L(\mathbf{w})$, gradient $\nabla L(\mathbf{w})$ points towards in direction of steepest ascent.


In 1d, either points left or right

Algorithm:

Take derivative

Move slightly in other
direction

Repeat



Gradient descent

Based on Joseph Redmon's slides,
UW Computer Vision class

For some loss function $L(\mathbf{w})$, gradient $\nabla L(\mathbf{w})$ points towards in direction of steepest ascent.


In 1d, either points left or right

Algorithm:

Take derivative

Move slightly in other
direction

Repeat



Gradient descent

Based on Joseph Redmon's slides,
UW Computer Vision class


Algorithm:

Take derivative

Move slightly in other
direction

Repeat

End up at local optima




Gradient descent

Based on Joseph Redmon's slides,
UW Computer Vision class

Formally:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \nabla L(\mathbf{w})$$

Where η is *step size*, how far to step relative to the gradient



Optimization in Classification

- **Decision Trees**



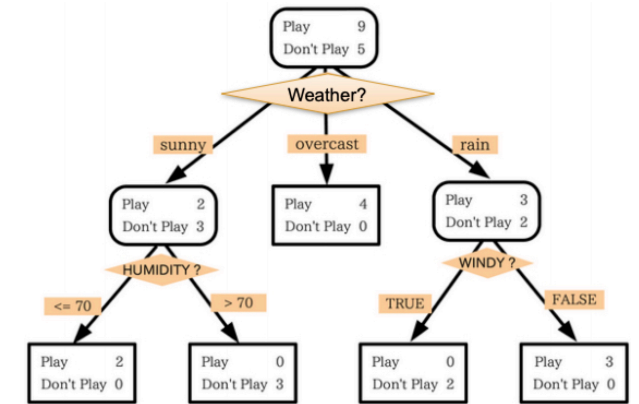
- Very hard to find optimal tree (min. misclassification)
- Top-down heuristic as shown in Lecture 1

- **Support Vector Machines**

- Convex optimization!



- **Logistic Regression**

- Convex optimization!



Recap

- Regression models:
 - Linear Regression
 - Ridge Regression
 - Lasso
- Overfitting and Underfitting
- Bias-Variance Tradeoff
- Gradient Descent for ML



Error = Noise + Bias + Variance

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \nabla L(\mathbf{w})$$